Received: 21 May 2016 | Accepted: 1 December 2016

DOI: 10.1111/faf.12200

ORIGINAL ARTICLE

WILEY FISH and FISHERIES

Improving estimates of population status and trend with superensemble models

Sean C Anderson[1] | Andrew B Cooper[1] | Olaf P Jensen[2] | Cóilín Minto[3] | James T Thorson[4] | Jessica C Walsh[1] | Jamie Afflerbach[5] | Mark Dickey-Collas[6,7] | Kristin M Kleisner[8] | Catherine Longo[9] | Giacomo Chato Osio[10] | Daniel Ovando[11] | Iago Mosqueira[10] | Andrew A Rosenberg[12] | Elizabeth R Selig[13]

[1]School of Resource and Environmental Management, Simon Fraser University, Burnaby, BC, Canada

[2]Institute of Marine & Coastal Sciences, Rutgers University, New Brunswick, NJ, USA

[3]Marine and Freshwater Research Centre, Galway-Mayo Institute of Technology, Galway, Ireland

[4]Fisheries Resource Analysis and Monitoring Division, Northwest Fisheries Science Center, National Marine Fisheries Service, National Oceanographic and Atmospheric Administration, Seattle, WA, USA

[5]National Center for Ecological Analysis and Synthesis, University of California Santa Barbara, Santa Barbara, CA, USA

[6]International Council for the Exploration of the Sea, Copenhagen, Denmark

[7]DTU Aqua National Institute of Aquatic Resources, Technical University of Denmark (DTU), Charlottenlund, Denmark

[8]Ecosystem Assessment Program, Northeast Fisheries Science Center, National Marine Fisheries Service, National Oceanographic and Atmospheric Administration, Woods Hole, MA, USA

[9]Marine Stewarship Council, London, UK

[10]Unit D.02 Water and Marine Resources, Directorate D – Sustainable Resources, DG Joint Research Centre, European Commission, Ispra, Italy

[11]Bren School of Environmental Science and Management, University of California Santa Barbara, Santa Barbara, CA, USA

[12]Union of Concerned Scientists, Cambridge, MA, USA

[13]Conservation International, Arlington, VA, USA

Correspondence
Sean C Anderson, School of Aquatic and Fishery Sciences, University of Washington, Seattle, WA, USA.
Email: sean.anderson@dal.ca

Funding information
Gordon and Betty Moore Foundation

Abstract

Fishery managers must often reconcile conflicting estimates of population status and trend. Superensemble models, commonly used in climate and weather forecasting, may provide an effective solution. This approach uses predictions from multiple models as covariates in an additional "superensemble" model fitted to known data. We evaluated the potential for ensemble averages and superensemble models (ensemble methods) to improve estimates of population status and trend for fisheries. We fit four widely applicable data-limited models that estimate stock biomass relative to equilibrium biomass at maximum sustainable yield (B/B_{MSY}). We combined these estimates of recent fishery status and trends in B/B_{MSY} with four ensemble methods: an ensemble average and three superensembles (a linear model, a random forest and a boosted regression tree). We trained our superensembles on 5,760 simulated stocks and tested them with cross-validation and against a global database of 249 stock assessments. Ensemble methods substantially improved estimates of population status and trend. Random forest and boosted regression trees performed the best at estimating population status: inaccuracy (median absolute proportional error) decreased from 0.42 – 0.56 to 0.32 – 0.33, rank-order correlation between predicted and true status

improved from 0.02 – 0.32 to 0.44 – 0.48 and bias (median proportional error) declined from −0.22 – 0.31 to −0.12 – 0.03. We found similar improvements when predicting trend and when applying the simulation-trained superensembles to catch data for global fish stocks. Superensembles can optimally leverage multiple model predictions; however, they must be tested, formed from a diverse set of accurate models and built on a data set representative of the populations to which they are applied.

KEYWORDS
CMSY, data-limited fisheries, ensemble methods, multimodel averaging, population dynamics, sustainable resource management

1 | INTRODUCTION

Status and trend are two of the most fundamental values to quantify in the management of ecological populations (e.g. Hutchings, Minto, Ricard, Baum, & Jensen, 2010; IUCN 2015). However, managers are often faced with reconciling multiple uncertain and potentially conflicting estimates of status and trend (e.g. Branch, Jensen, Ricard, Ye, & Hilborn, 2011; Brodziak & Piner, 2010; Deroba et al., 2015). For example, one model may suggest a population is at risk and declining in abundance while others may suggest it is not at risk and stable.

One solution is to take the average or weighted average of several model predictions, that is, an ensemble. Such ensembles are typically more accurate and less biased than individual model estimates and can incorporate various types of uncertainty, such as uncertainty in model structure, initial conditions and parameter estimation (Araújo & New, 2007; Dietterich, 2000). Ensembles are superior to individual models in at least three ways: (i) statistically by averaging across models and therefore being less likely to pick the "wrong" model, (ii) computationally by reducing the risk of getting stuck in local optima, and (iii) representationally by expanding the range of hypotheses explored (Dietterich, 2000). This approach forms the basis of many machine-learning methods (e.g. Dietterich, 2000), has helped reconcile climate forecasts from dozens of models (e.g. IPCC 2013; Murphy et al., 2004; Tebaldi & Knutti, 2007) and even improved early warning signs of malaria outbreaks (Thomson et al., 2006). In ecology, ensemble methods are sometimes used to improve species distribution modelling (e.g. Araújo & New, 2007; Breiner, Guisan, Bergamini, & Nobis, 2015) and indeed have been used to combine estimates of population status and trend (e.g. Brodziak & Piner, 2010).

Whereas averages or weighted averages of model estimates may improve predictions compared to a single model, they may not optimally leverage available data. The best prediction does not necessarily lie in the middle of multiple model predictions, some models may perform better than others in certain conditions, and the covariance between models may contain information that can improve predictive accuracy. For example, one model might perform well at estimating high levels of abundance but be biased at low levels of abundance, while another model might have the opposite properties. An optimal combination of these models is not simply an average of the two.

We can exploit these characteristics by using the predictions from a group of models as inputs into a separate statistical model. This technique, sometimes called superensemble modelling (Krishnamurti et al., 1999), is common in climate and weather forecasting (e.g. Mote et al., 2015; Yun et al., 2005). The superensemble is fit to a training data set where outcomes are well known and then used to predict on a data set of interest. For example, Krishnamurti et al. (1999) combined predictions of wind and precipitation in Asian monsoons via a superensemble regression fit to observed data. Their superensemble was considerably more accurate than any individual prediction or an average of the predictions.

In fisheries science, the commonly used operational models for determining status and trend of exploited fish populations are stock assessments; that is, population models coupled to an observation model that incorporate all appropriate data (e.g. catches, size and age distributions, surveys, and tagging information) to quantify values such as the biomass of a stock that can produce maximum sustainable yield (B_{MSY}; Hilborn & Walters, 1992). However, the broad range of data required to conduct these stock assessments is not available for the majority of fish populations, including those of conservation concern and of economic interest to fisheries (FAO 2014). Therefore, a number of models have been proposed to assess B/B_{MSY} based on the

limited data available for the majority of fish stocks: (i) a time series of the total weight of catch, and (ii) a basic understanding of population productivity (e.g. Martell & Froese, 2013; Vasconcellos & Cochrane, 2005). Recently, Rosenberg et al. (2014) investigated the performance of four data-limited models through a large-scale simulation experiment. Three of these models were based on Schaefer (logistic) biomass dynamics, and one was an empirical model fitted to more data-rich stock-assessment output. The four models frequently disagreed about population status (e.g. Figure 1), no one model had strong performance across all fish stocks, and some models performed better than others depending on circumstances.

Here, we estimate population status and trend of exploited fish populations using ensembles and superensembles (collectively "ensemble methods") of these four data-limited models. We apply four ensemble-method approaches of varying complexity to both simulated and real-world fish stocks and compare their predictive performance against each other and the individual models.

2 | METHODS

To test the ability of superensembles to improve estimates of status and trend in data-limited fish stocks, we first fit four individual assessment models to a large simulated data set of fish stocks. We then built and tested the performance of superensembles using cross-validation of the simulated data set. Finally, we tested superensembles built with the entire simulated data set against a database of global fish stocks. We describe these steps in detail below and illustrate the general approach graphically and with pseudocode in Figure 2.

2.1 | Individual models of population status

We fit four individual data-limited models that use catch data and basic life-history parameters to estimate B/B_{MSY}. We chose these models because they can be fit to the vast majority of fisheries around the world, are established in the literature and have been extensively simulation tested (Rosenberg et al., 2014).

Three of the models are mechanistic and based generally on Schaefer biomass dynamics (Schaefer, 1954) of the form

$$\hat{B}_{t+1} = B_t + rB_t(1 - B_t/B_0) - C_t,$$

where $\hat{B}_{t+1}$ represents predicted biomass at time t plus 1 year, B_t represents biomass at time t, r represents intrinsic population growth rate, B_0 represents unfished biomass or carrying capacity K, and C represents catch. The fourth model is an empirically derived model based on the RAM Legacy Stock Assessment Database (Ricard, Minto, Jensen, & Baum, 2012). Rosenberg et al. (2014) provide a full background on these four methods (http://www.fao.org/docrep/019/i3491e/i3491e00.htm, last accessed 2016-11-08), and code to fit all the models is available in an accompanying package datalimited for the statistical software R (R Core Team 2015) https://github.com/datalimited/datalimited (last accessed 2016-11-08) or Anderson et al. (2016a). In summary:

- CMSY (catch-MSY) implements a stock-reduction analysis with Schaefer biomass dynamics (Martell & Froese, 2013). It requires a prior distribution on r and K as well as priors on the relative proportion of biomass at the beginning and end of the time series compared to unfished biomass (depletion). The version of the model used in Rosenberg et al. (2014) was modified from Martell and Froese (2013) to generate biomass trends from all viable r-K pairs and produce an estimate of B/B_{MSY} from the median trend.
- COM-SIR (catch-only model with sampling-importance resampling) is a coupled harvest-dynamics model (Vasconcellos & Cochrane, 2005). Biomass is assumed to follow a Schaefer model and harvest dynamics are assumed to follow a logistic model. The model is fit with a sampling-importance-resampling algorithm (Rosenberg et al., 2014).
- SSCOM (state-space catch-only model) is a hierarchical model that, similar to COM-SIR, is based on a coupled harvest-dynamics model (Thorson, Minto, Minte-Vera, Kleisner, & Longo, 2013). SSCOM estimates unobserved dynamics in both fishing effort and the fished population based on a catch time series and priors on r, the maximum rate of increase of fishing effort and the magnitude of various forms of stochasticity. The model is fit in a Bayesian state-space framework to integrate across three forms of stochasticity: variation in effort, population dynamics and fishing efficiency (Thorson et al., 2013).
- mPRM (modified panel-regression model) is a modified version of the panel-regression model from Costello et al. (2012). Unlike the other models, mPRM is empirical and not mechanistic—it uses the RAM Legacy Stock Assessment Database to fit a regression model to a series of characteristics of the catch time series and stock with stock-assessed B/B_{MSY} as the response. The model used in this article is modified from the original—it condenses the life-history categories into three categories to match the simulated data set, removes the maximum catch predictor because the absolute catch in the simulated data set is arbitrary, and does not implement the bias correction needed in Costello et al. (2012) because we do not derive estimates of median status across multiple stocks.

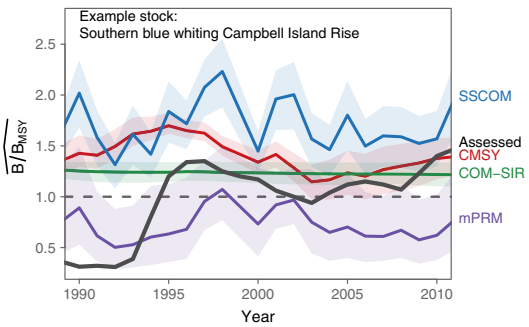

FIGURE 1 Different models can suggest conflicting population statuses and trends. Shown are trajectories of estimated B/B_{MSY} from four data-limited assessment methods (colours) and a data-rich stock assessment (black) for Southern blue whiting (*Micromesistius australis*) on Campbell Island Rise, New Zealand. Lines indicate median fits, and shaded regions indicate interquartile ranges. Dashed horizontal line indicates $B/B_{MSY} = 1$

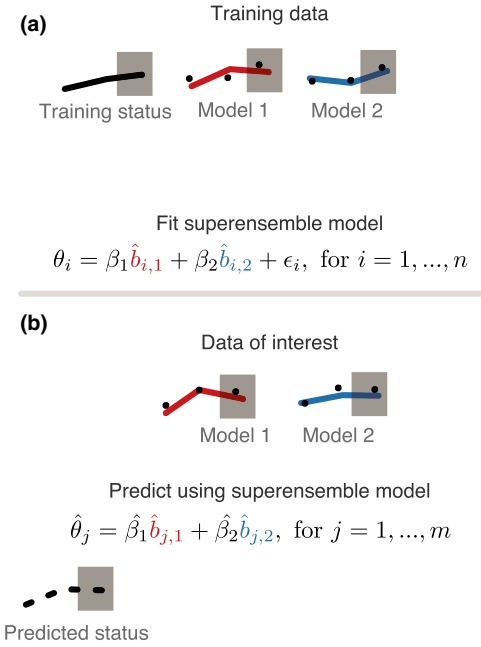


```
# Fit models to training data:
cmsy_status <- cmsy(...)
comsir_status <- comsir(...)
# Combine and reshape output into wide format (not shown)
head(data_training)
# > cmsy_status comsir_status known_status
# > 1.4         1.3           1.2
# > ...         ...           ...
# Build superensemble model:
ensemble_model <- lm(known_status ~ cmsy_status +
  comsir_status, data = data_known)
```

```
# Fit models to data of interest:
cmsy_status <- cmsy(...)
comsir_status <- comsir(...)
# Combine and reshape output into wide format (not shown)
head(data_interest)
# > cmsy_status comsir_status ...
# > 0.9         1.1
# > ...         ...
# Predict status using superensemble model:
predict(ensemble_model, newdata = data_interest)
# > 1.06  0.98  0.85 ... # status across multiple fisheries
```

FIGURE 2 Using a superensemble model to predict population status from two individual models. The process is illustrated graphically on the left and with R pseudocode on the right. (a) Individual models (red and blue lines) are fit to training data (dots) from populations of known or assumed status (known status shown by black line). The shaded grey boxes indicate the recent time period that we are interested in for this article. Estimates of status from these individual models ($\hat{b}_{i,1}$ and $\hat{b}_{i,2}$), potentially combined with additional covariates, are then used as covariates in a statistical model fitted to the known or assumed population status as the response (here represented as a linear model). The symbols β and ε represent parameters and error in the linear model, respectively. The *i* subscripts represent individual fish stocks from 1 to *n*, and θ represents the known status. (b) The superensemble can then be used to make predictions for new stocks of interest. The same individual models are fit to populations of interest and then combined using the previously fitted superensemble model. Here, the *j* subscripts represent individual fish stocks from 1 to *m*, and $\hat{\theta}$ represents the predicted status. The $\hat{\beta}$ represents the parameters estimated when the superensemble was fit in panel (a)

2.2 | Simulated data set to build the superensemble

We first developed and tested ensemble methods on a fully factorial simulated data set of fisheries with known status (Rosenberg et al., 2014). Briefly, these simulations were implemented with the FLR packages (Kell et al., 2007) for the statistical software R, and, in particular, the FLBRP package. The framework takes a series of life-history parameters and fishery characteristics to generate population projections and resulting catch time series. Life-history values (e.g. mean asymptotic length) for three fish life histories (small pelagic, demersal and large pelagic) were translated into a complete set of parameters for a von Bertalanffy growth model, a maturity ogive, natural mortality, a selectivity function and a Beverton–Holt stock-recruitment function using the life-history relationships derived in Gislason, Pope, Rice, and Daan (2008).

Fishing scenarios included three levels of initial biomass depletion compared to carrying capacity: biomass at 100%, 70% and 40% of carrying capacity; and four exploitation patterns: (i) a constant exploitation rate, (ii) an exploitation rate coupled with biomass to mimic an open-access single-species fishery, (iii) a scenario where exploitation rate increased continuously, and (iv) a "roller coaster" scenario where the exploitation rate increased and then decreased. Process noise (recruitment variability; that is, unexplained variability in population dynamics) was introduced to the models at two magnitudes in log space, N (0,0.2^2) and *N* (0,0.6^2), and was either uncorrelated through time or had a first-order autoregressive parameter of .6. The simulation also included a scenario with N (0,0.2^2) measurement error around log catch and one scenario without measurement error. Rosenberg et al. (2014) ran 10 iterations for each combination of factors adding stochastic draws of recruitment and catch-recording variability each time to generate a total of 5,760 stocks. Code to generate the simulations is available at https://github.com/datalimited/stocksims (last accessed 2016-11-08).

2.3 | Building the superensemble models

The individual models we seek to combine with superensembles provide time series of stock status (B/B_{MSY}). Therefore, we can use superensembles to estimate any property of these time series. Here, we focus on two properties: the mean and slope of B/B_{MSY} in the last 5 years. Together, these quantities address the recent state and trend of stock status, which are both of management and conservation interest (e.g. Hutchings et al., 2010; IUCN 2015). To avoid undue influence of the time series end points on the calculated slope, we measured the slope as the Theil–Sen estimator of median slope (Theil, 1950).

We used the mean or slope of B/B_{MSY} as the response variable and the predictions from the individual models as predictors in our superensemble models (Figure 2a). When modelling mean B/B_{MSY}—a ratio bounded at zero—we fit the superensemble models in log space

and exponentiated the predictions. For the estimates of B/B_{MSY} slope, which are not bounded at zero, we fit superensemble models on the natural untransformed scale.

We compared an ensemble average and three superensembles of varying complexity: a linear model with two-way interactions, a random forest and a boosted regression tree. We describe these models as estimating $\hat{\theta}$, which represents either the ensemble estimated log B/B_{MSY} or slope of B/B_{MSY}. The individual model estimates of log B/B_{MSY} or slope of B/B_{MSY} are represented as $\hat{b}$ for models 1 through 4 (CMSY, COM-SIR, SSCOM, mPRM). The ensemble average for each fishery i was calculated as:

$$\hat{\theta}_i = (\hat{b}_{i,1} + \hat{b}_{i,2} + \hat{b}_{i,3} + \hat{b}_{i,4})/4, \text{for } i = 1, \ldots, n.$$

We fit the linear model superensemble with all second-order interactions:

$$\hat{\theta}_i = \beta_0 + \beta_1 \hat{b}_{i,1} + \ldots + \beta_{1,2} \hat{b}_{i,1} \hat{b}_{i,2} + \ldots + \epsilon_i, \epsilon_i \sim \mathrm{N}(0, \sigma^2), \text{for } i = 1, \ldots, n.$$

For this illustrative example, we chose this level of model complexity a priori but a modeller could apply model selection via information–theoretic or cross-validation approaches.

Our two machine-learning superensemble models, a random forest and a generalized boosted model (GBM), were based on regression trees. Regression trees sequentially determine what value of a predictor best splits the response data into two branches based on a loss function (Breiman, Friedman, Stone, & Olshen, 1984). In random forests, a series of regression trees are built on a random subset of the data and random subset of the covariates of the model (Breiman, 2001). In GBMs, each subsequent model is fit to the residuals from the previous model; data points that are fit poorly in a given model are given more weight in the next model (Elith, Leathwick, & Hastie, 2008). Random forests and GBMs can provide strong predictive performance and fit highly nonlinear relationships (Elith et al., 2008; Hastie, Tibshirani, & Friedman, 2009). We fit random forest models with the randomForest package (Liaw & Wiener, 2002) for R with the default argument values. We fit boosted regression tree models with the gbm package (Ridgeway, 2015) for R. We fit GBMs with 2,000 trees, an interaction depth of 6, a learning rate (shrinkage parameter) of 0.01 and all other arguments at their default values.

2.4 | Additional covariates

Superensemble models allow us to incorporate additional covariates and potentially leverage interactions between these covariates and individual model predictions. Additional covariates could be, for example, life-history characteristics, information on exploitation patterns or statistical properties of the data. We tested the performance benefits of including one set of additional covariates: spectral properties of the catch time series. Spectral analysis decomposes a time series into the frequency domain and provides a means of describing the cyclical shape of the catch series that is independent of time-series length (except in affecting precision) and independent of absolute magnitude of catch. We fit spectral models to the scaled catch time series (catch divided by maximum catch) with the spec.ar function in R and recorded representative short- and long-term spectral densities at frequencies of 0.20 and 0.05, which correspond to 5- and 20-year cycles. For the linear model superensemble, we incorporated the two spectral covariates ($S1$, $S2$) along with all second-order interactions as:

$$\hat{\theta}_i = \beta_0 + \beta_1 \hat{b}_{i,1} + \ldots + \beta_{1,2} \hat{b}_{i,1} \hat{b}_{i,2} + \ldots + \beta_{S1} S1_i + \beta_{S2} S2_i + \beta_{S1,S2} S1_i S2_i + \beta_{1,S2} \hat{b}_{i,1} S2_i + \ldots + \epsilon_i,$$

with $\varepsilon \sim \mathrm{N}(0, \sigma^2)$ and for simulated fisheries i 1 through n. We include the results of adding these additional covariates in the supplementary materials.

2.5 | Applying the superensemble models and testing performance

Once the superensemble models are built and trained using the simulated stocks (or any data set with "known" status), we can use the superensembles to estimate the status of new stocks (Figure 2b). To do this, we applied the individual models to our stocks of interest (i.e. CMSY, COM-SIR, SSCOM, mPRM) and then used these individual model estimates of status or trend as data in our already built superensemble models. In this article, we applied the superensemble models to subsets of the simulated data as a cross-validation test to test predictive performance and to the RAM Legacy Stock Assessment Database to test predictive performance on real stocks.

We used repeated threefold cross-validation: we randomly divided the data set into three sets, built superensemble models on two-thirds of the data and evaluated predictive performance on the remaining third. We repeated this across each of the three splits and then repeated the whole procedure 50 times to account for bias that may result from any one set of validation splits. In the simulated data set, there were 10 replicates of each unique combination of simulation parameters that differed only in stochastic variability. As the dynamics of these populations were often similar, we grouped these stocks in the cross-validation process into either the training or testing split.

We also tested our ensemble methods on the RAM Legacy Stock Assessment Database—a compilation of stock-assessment output from hundreds of exploited marine populations around the world. Our analysis of the stock-assessment database was based on version 2.5. After removing stocks for which at least one of the individual models did not converge (121), this database included 249 stocks. We removed these stocks for all methods—both for the individual and superensemble models. An alternative would be to fit separate superensemble models to subsets of the individual models that did converge, but for simplicity, we only used superensemble models fitted to all four individual models.

In the case of the RAM Legacy Stock Assessment Database, we used superensembles trained on the entire simulation data set. However, because mPRM is built on the same stock-assessment database, we applied threefold cross-validation to the data underlying the mPRM model so that the data set with which mPRM was

trained (for the individual model and superensemble) was separate from the data set with which it was tested. This meant that, for each iteration of cross-validation, we split the RAM database into three, fit the mPRM model to two-thirds of the RAM database, fit a superensemble with this version of mPRM and then tested the performance of the superensemble on the third of the RAM database we had withheld.

Predictive performance can be evaluated with metrics that represent a variety of modelling goals. For continuous response variables such as the mean and slope of population status, performance metrics often measure some form of bias, precision, accuracy (a combination of bias and precision) or the ability to correctly rank or correlate across populations (e.g. Walther & Moore, 2005). Here, we measure proportional error, defined as $(\hat{\theta}-\theta)/|\theta|$, where $\hat{\theta}$ and θ represent estimated and "true" (or stock-assessed) mean or slope of B/B_{MSY}. We calculated median proportional error to measure bias, median absolute proportional error to measure accuracy and Spearman's rank-order correlation between predicted and "true" values to measure the ability to correctly rank populations. When testing with the RAM Legacy Stock Assessment Database, we treated the estimates from these data-rich stock assessments as known without error. Thus, any error in the stock-assessment estimates of the mean or slope of B/B_{MSY} also contributes to our estimates of prediction error for each of the four data-limited models and the ensembles. Code to reproduce our analysis is available at https://github.com/datalimited/ensembles (last accessed 2016-11-08) or Anderson et al. (2016b).

3 | RESULTS

Applied to the simulated data set of known stock status, the individual models had variable success at estimating the mean (status) and slope (trend) of B/B_{MSY} in the last 5 years. All models exhibited a high degree of scatter around the one-to-one line of perfect status prediction (Figure 3). In contrast to the known unimodal distribution of status, CMSY exhibited bimodal predictions (Figure 3a), but had the best rank-order correlation and accuracy scores (Figure 4a). COM-SIR and SSCOM both correctly identified a number of stocks with low status, but frequently predicted a high status when status was in fact low (Figure 3b,c). mPRM had relatively poor ability to predict status for the simulated data set (Figure 3d). There was generally little correlation between true and predicted recent trend in status for any of the individual models (rank-order correlation = .02–.25) with the exception of SSCOM (correlation = .54; Figure S1a–d).

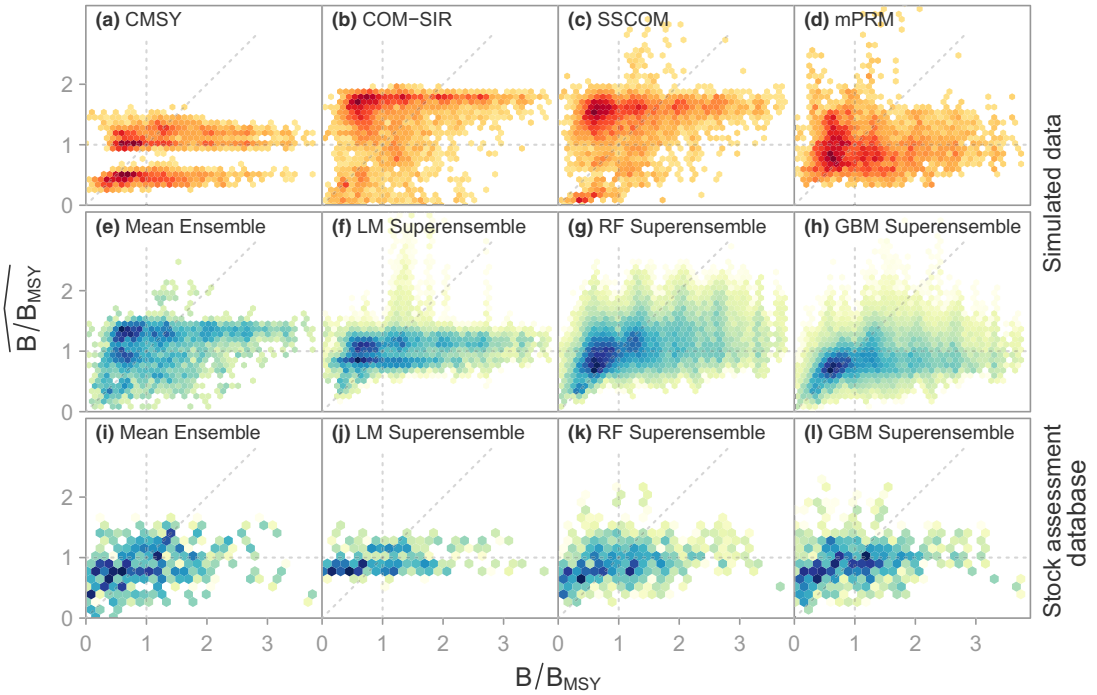

FIGURE 3 True (or assessed) population status (*x*-axis) vs predicted population status from individual models and ensemble methods with cross-validation (*y*-axis). These scatterplots represent the aggregate results of repeated threefold cross-validation tests where the ensemble models are built on two-thirds of the data and tested on the remaining third. (a–d) Individual data-limited model estimates of mean B/B_{MSY} (biomass divided by biomass at maximum sustainable yield) in the last 5 years for a simulated data set of known population status. (e–h) Ensemble estimates for the same populations. Shown are a mean ensemble, a linear superensemble model with two-way interactions (LM), a random forest superensemble (RF) and a generalized boosted regression model superensemble (GBM). (i–l) The same ensemble models, which were trained on the simulated data set, applied to the RAM Legacy Stock Assessment Database and compared to data-rich stock-assessed status. In the case of the RAM Legacy Stock Assessment data, we refit the modified panel-regression model (mPRM) on each cross-validation split. We binned the data into hexagons for visual presentation. Darker areas indicate areas with greater density of data. Yellow–red shading and yellow–blue shading distinguishes individual models from ensemble methods

Ensemble methods, and in particular the machine-learning superensemble models (random forest and GBM), generally improved estimates of status and trend over any individual model (Figure 3e–h, Figure S1e–h). Compared to the individual models, machine-learning superensembles decreased inaccuracy (median absolute proportional error) from 0.42 – 0.56 to 0.32 – 0.33, increased rank-order correlation from 0.02 – 0.32 to 0.44 – 0.48 and reduced bias (median proportional error) from −0.22 – 0.31 to −0.12 – 0.03 (Figure 4a). These superensembles also generally had better ability to distinguish if simulated stocks were above or below B/B_{MSY} = 0.5 (Figure S2). Results were similar when predicting trend: compared to individual models, machine-learning superensembles decreased inaccuracy from 0.04 – 0.06 to 0.03 – 0.03, increased rank-order correlation from 0.02 – 0.54 to 0.61 – 0.65 and reduced bias from −0.009 – 0.014 to −0.002 – 0.002 (Figure S3). The ensemble models that simply took a mean of the individual models ranked slightly behind the best individual model for estimating fish stock status (CMSY; Figure 4a) and had slightly lower correlation but higher accuracy than the best individual model at predicting the trends of status (SSCOM; Figure S3).

The superensemble models were able to improve the predictive performance by harnessing the best properties of individual models, the covariance between individual models, and interactions with other covariates. For example, SSCOM had strong predictive ability when it predicted low B/B_{MSY} (Figure 3c, Figure S4c) and CMSY predictions were approximately linearly related to B/B_{MSY} within the low and high clusters of predictions (Figure S4). SSCOM contributed most strongly on its own to determining trend (Figure S5). Superensembles also exploited the covariance between individual model predictions. For instance, both the linear model and GBM ensemble suggest that if mPRM and SSCOM predict high status, the true status also tends to be high (Figures S6 and S7f). The addition of spectral density covariates helped the superensemble models correctly predict higher status values (Figure S8g,h). The performance of the ensembles was only marginally improved by including these covariates (Figure S9 vs Figure 4).

When applied to the stock-assessment database, the superensemble models—trained exclusively on the simulated data set—generally performed as well or better than the best individual models. The mean, random forest and GBM ensembles outperformed the mPRM method which is trained directly on the RAM Legacy Stock Assessment Database itself (Figure 4b, Figure S10). Compared to the individual models, the machine-learning superensembles increased accuracy by 0%–30%, improved correlation from 0.19 – 0.36 to 0.35 – 0.38 and reduced bias from −0.25 – 0.45 to −0.05 – 0.02.

4 | DISCUSSION

Ensemble methods provide a useful approach to situations where environmental resource management decisions must be made on the basis of multiple, potentially contrasting estimates of status. Compared to individual models of fish population status, ensemble methods were consistently the best or among the best across three performance dimensions (accuracy, bias and rank-order correlation), two response variables (status and trend), two data sets (simulated and global fisheries) and multiple ensemble methods (from a simple average to machine-learning superensembles). Our results suggest choosing a superensemble model that allows for nonlinear relationships, such as machine-learning methods; these models provided added insight into individual model behaviour and generally performed the best.

Certain conditions will make some ensemble models more effective than others. First, ensembles will be most effective when they are comprised of diverse individual models that choose different structural model forms, explore contrasting but plausible ranges of parameter values and make uncorrelated errors (Ali & Pazzani, 1996; Dietterich, 2000; Tebaldi & Knutti, 2007). We would expect such models to perform well in different conditions and an ensemble model can exploit the best predictive performance of each. Second, ensemble models will be most effective when they are not overfit to the training data set. Cross-validation testing (Caruana, Niculescu-Mizil, Crew, &

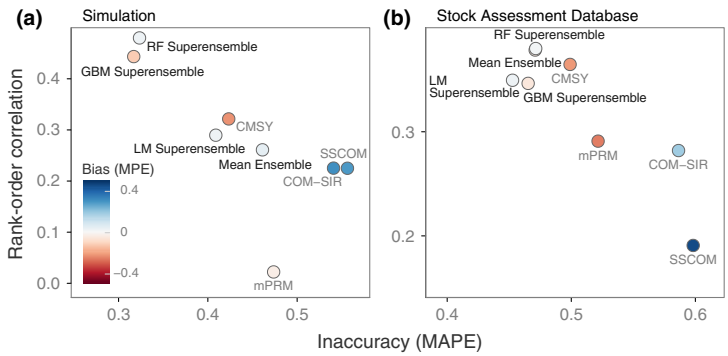

FIGURE 4 Performance metrics of individual and ensemble models predicting mean B/B_{MSY} (biomass divided by biomass at maximum sustainable yield) in the last 5 years fitted to a data set with (a) known population status and (b) the RAM Legacy Stock Assessment Database. The *x*-axis represents within-population inaccuracy: median absolute proportional error (MAPE). The *y*-axis represents across-population Spearman rank-order correlation. The top-left corner contains methods with the best performance across the two metrics. The colour shading represents bias (median proportional error; MPE): white points are unbiased, blue points represent methods that predict B/B_{MSY} values that are too high, red points represent methods that predict B/B_{MSY} values that are too low. These performance metrics are derived from the data in Figure 3 and based on repeated threefold cross-validation testing

Ksikes, 2004; Hastie et al., 2009) and methods that are robust to overfitting such as random forests (Breiman, 2001) may help avoid overfitting ensemble models. We note that our simplest ensemble model, an average of individual model predictions, performed approximately as well as complex machine-learning models when we trained our superensembles on the simulation data set and tested them on a separate "real" data set (i.e. the RAM Legacy Stock Assessment Database, Figure 4b). Third, ensemble models will be most effective when they are trained on data that are representative of the data set of interest (Knutti, Furrer, Tebaldi, Cermak, & Meehl, 2009; Weigel, Knutti, Liniger, & Appenzeller, 2010). Cross-validation within a training data set will provide an optimistically biased impression of predictive performance if the training data set fundamentally differs from the data set of interest (Hastie et al., 2009).

We illustrated that superensembles can improve point estimates of population status and trends in status; however, there is no reason why superensembles cannot also be used to provide measures of uncertainty around those point estimates. The same approaches to deriving measures of uncertainty from any regression model are available to a superensemble. For example, likelihood profile confidence intervals or Bayesian credible intervals are available for superensembles fit via maximum likelihood or Bayesian procedures, respectively. Measures of predictive uncertainty can be generated for machine-learning methods such as random forests or GBMs using bootstrap procedures (e.g. Finnegan et al., 2015; Hastie et al., 2009). Furthermore, uncertainty from the component models could be included in superensembles. These superensembles could be fit using any errors-in-variables or measurement-error modelling approach (e.g. Carroll, Ruppert, Stefanski, & Crainiceanu, 2006).

Multimodel inference in the form of coefficient averaging weighted by information theoretics such as the Akaike information criterion (AIC) is a common analytical approach in fisheries and ecology (e.g. Burnham & Anderson, 2002; Grueber, Nakagawa, Laws, & Jamieson, 2011; Johnson & Omland, 2004). The ensemble methods described in this article share similarities with coefficient averaging but differ in other important ways. Ensemble methods and coefficient averaging share the long-held notion that multiple working hypotheses can contribute useful information for inference (Chamberlin, 1890). A fundamental difference is that coefficient averaging focuses on averaging coefficients whereas ensembles instead average predictions. Thus, ensembles provide a general purpose tool: they do not require information theoretics, and they can combine different types of models (e.g. parametric and nonparametric models or frequentist and Bayesian predictions). Furthermore, superensembles extend these benefits by allowing model predictions to be combined via nonlinear functions that are tuned to known data.

A strength of superensembles is that they can be tailored to predict specific response variables. For example, we built separate superensemble models of mean B/B_{MSY} and the slope of B/B_{MSY}. The same set of model weights or nonlinear relationships need not hold across different response variables. For instance, SSCOM contributed little to the GBM superensemble estimate of status at higher levels of predicted B/B_{MSY} (Figure S4), but contributed strongly to estimates of trend (Figure S5). Formally, fitting superensemble models to specific quantities of interest (such as the slope of B/B_{MSY}) provides an additional calibration step to a quantity of interest (Rykiel, 1996). This ensemble calibration could include a loss function tailored to the goals of the model, say placing greater weight on accuracy at lower rather than higher status levels. Conversely, because superensembles are tailored to a specific response and loss function, superensembles force a modeller to choose an operational purpose for their model upfront (*sensu* Dickey-Collas, Payne, Trenkel, & Nash, 2014). For instance, one could have an ensemble estimate of B and an ensemble estimate of B_0, but their ratio may not be the same as an ensemble estimate of B/B_0. A modeller might therefore choose to focus on B/B_0, which provides a unitless ratio, is easier to compare across stocks, and the ratio is often a more stable estimate across models (Deroba et al., 2015).

As Box and Draper (1987) noted, all models are wrong, but some may still be useful. The ensemble methods we investigated attempt to piece together the useful parts of candidate models to build a model with improved performance. Instead of viewing the superensemble as a black box, we think considerable mechanistic understanding can be gained by studying its structure. For example, when SSCOM estimates low status this is likely the case, conversely when COMSIR estimates low status, the true status is more likely to be high (Figure S4). These models have two main differences: (i) the form of effort dynamics, and (ii) the allowance for both measurement and process error in SSCOM, whereas the implemented COMSIR admits measurement error only. Were the methods to differ only in effort dynamics, the results point towards a more suitable representation of effort dynamics at low biomasses in SSCOM. We think that such investigation of the structure of a superensemble may lead to improvement in the mechanisms assumed in individual models.

Combining predictions from multiple models via superensemble methods is broadly useful in other subfields of fisheries science and ecology in general. In fisheries science, superensembles provide an additional tool to assist with some longstanding issues. For example, superensembles are helpful because modellers need not decide on one model—instead of deciding on dome versus asymptotic fisheries selectivity (e.g. Sampson & Scott, 2012), or on whether to fix or estimate natural mortality (e.g. Johnson et al., 2015), superensembles can use multiple models to draw inference. Furthermore, the relative contributions of individual models can help tease apart the conditions under which various model assumptions result in the most accurate predictions. Finally, superensembles can be used to directly estimate other quantities of interest in fisheries science. For instance, superensembles could help assess overfishing by estimating fishing mortality compared to fishing mortality at MSY (F/F_{MSY}) or be trained to estimate natural mortality.

More broadly, in ecology, predictions about extinction risk are widely used at national (e.g. the US Endangered Species Act and the Canadian Species at Risk Act) and international (e.g. the IUCN Red List, IUCN 2015) levels. These risk assessments generally involve fitting regression models to outcomes for individual species along with predictors of extinction risk (e.g. Anderson, Farmer, Ferretti, Houde, & Hutchings, 2011; Pinsky, Jensen, Ricard, & Palumbi, 2011), or fitting

population-dynamic models to data for individual species (e.g. DFO 2010). Both types of models are prone to error caused by model misspecification and therefore results are sensitive to decisions about model structure (Brooks & Deroba, 2015). Although there are options to account for potential model misspecification in determination of species risk (e.g. coefficient averaging, Burnham & Anderson, 2002; generalized modelling, Yeakel, Stiefs, Novak, & Gross, 2011; or semi-parametric methods, Thorson, Ono, & Munch, 2014), ensemble methods are a relatively simple way to combine predictions in a transparent manner. Beyond estimates of status and trend, ensemble methods could be used, for example, to increase the robustness of spatial predictions when designing networks of protected areas (Rassweiler, Costello, Hilborn, & Siegel, 2014) or to forecast potential spatial shifts in species distribution given climate impacts (Harsch, Zhou, HilleRisLambers, & Kot, 2014). In any case, superensembles are not a panacea and are ultimately limited by the quality, breadth and representativeness of simulated or trusted data to which they are calibrated.

ACKNOWLEDGEMENTS

We thank members of Phase I of the working group "Developing new approaches to global stock status assessment and fishery production potential of the seas" who contributed to developing the data-limited methods and simulations used in our analysis. We thank E. Jardim, F. Scott and J.A. Hutchings for helpful comments during the development of this project and R.D. Methot for comments on an earlier version of the manuscript. We also thank two anonymous reviewers for their thoughtful reviews, comments and criticisms. We thank the Gordon and Betty Moore Foundation for funding the working group "Applying data-limited stock status models and developing management guidance for unassessed fish stocks."

REFERENCES

Ali, K. M., & Pazzani, M. J. (1996). Error reduction through learning multiple descriptions. *Machine Learning*, *24*, 173–202.

Anderson, S. C., Cooper, A. B., Jensen, O. P., Minto, C., Thorson, J. T., Walsh, J. C., … Selig, E. R. (2016a). *datalimited: Stock assessment methods for data-limited fisheries. R package version 0.1.0*. Geneva, Switzerland: Zenodo. doi:10.5281/zenodo.220991

Anderson, S. C., Cooper, A. B., Jensen, O. P., Minto, C., Thorson, J. T., Walsh, J. C., … Selig, E. R. (2016b). *Data and code for "Improving estimates of population status and trend with superensemble models"*. Geneva, Switzerland: Zenodo. doi:10.5281/zenodo.220990

Anderson, S. C., Farmer, R. G., Ferretti, F., Houde, A. L. S., & Hutchings, J. A. (2011). Correlates of vertebrate extinction risk in Canada. *BioScience*, *61*, 538–549.

Araújo, M. B., & New, M. (2007). Ensemble forecasting of species distributions. *Trends in Ecology & Evolution*, *22*, 42–47.

Box, G. E., & Draper, N. R. (1987). *Empirical model-building and response surfaces*. New York, NY: Wiley.

Branch, T. A., Jensen, O. P., Ricard, D., Ye, Y., & Hilborn, R. (2011). Contrasting global trends in marine fishery status obtained from catches and from stock assessments. *Conservation Biology*, *25*, 777–786.

Breiman, L. (2001). Random forests. *Machine Learning*, *45*, 5–32.

Breiman, L., Friedman, J., Stone, C. J., & Olshen, R. A. (1984). *Classification and regression trees*. Belmont, CA: Wadsworth International Group.

Breiner, F. T., Guisan, A., Bergamini, A., & Nobis, M. P. (2015). Overcoming limitations of modelling rare species by using ensembles of small models. *Methods in Ecology and Evolution*, *6*, 1210–1218.

Brodziak, J., & Piner, K. (2010). Model averaging and probable status of North Pacific striped marlin, *Tetrapturus audax*. *Canadian Journal of Fisheries and Aquatic Sciences*, *67*, 793–805.

Brooks, E. N., & Deroba, J. J. (2015). When "data" are not data: The pitfalls of post hoc analyses that use stock assessment model output. *Canadian Journal of Fisheries and Aquatic Sciences*, *72*, 634–641.

Burnham, K. P., & Anderson, D. R. (2002). *Model selection and multimodel inference: A practical information-theoretic approach*, 2nd edn. New York, NY: Springer.

Carroll, R. J., Ruppert, D., Stefanski, L. A., & Crainiceanu, C. M. (2006). *Measurement error in nonlinear models: A modern perspective*. Boca Raton, FL: Chapman & Hall/CRC.

Caruana, R., Niculescu-Mizil, A., Crew, G., & Ksikes, A. (2004). Ensemble selection from libraries of models. In: *Proceedings of the Twenty-first International Conference on Machine Learning (ICML)*. New York, NY.

Chamberlin, T. C. (1890). The method of multiple working hypotheses. *Science*, *15*, 92–96.

Costello, C., Ovando, D., Hilborn, R., Gaines, S. D., Deschenes, O., & Lester, S. E. (2012). Status and solutions for the world's unassessed fisheries. *Science*, *338*, 517–520.

Deroba, J. J., Butterworth, D. S., Methot, R. D., De Oliveira, J. A. A., Fernandez, C., Nielsen, A., … Hulson, P.-J. F. (2015). Simulation testing the robustness of stock assessment models to error: Some results from the ICES strategic initiative on stock assessment methods. *ICES Journal of Marine Science*, *72*, 19–30.

DFO (2010). Stock assessment update for British Columbia canary rockfish. Technical report, Canadian Science Advisory Secretariat, Ottawa, Canada.

Dickey-Collas, M., Payne, M. R., Trenkel, V. M., & Nash, R. D. M. (2014). Hazard warning: Model misuse ahead. *ICES Journal of Marine Science*, *71*, 2300–2306.

Dietterich, T. G. (2000). Ensemble methods in machine learning. In *Multiple classifier systems* (pp. 1–15). Berlin, Heidelberg: Springer.

Elith, J., Leathwick, J. R., & Hastie, T. (2008). A working guide to boosted regression trees. *Journal of Animal Ecology*, *77*, 802–813.

FAO (2014). The state of the world fisheries and aquaculture. Technical report, Food and Agriculture Organization of the United Nations, Rome.

Finnegan, S., Anderson, S. C., Harnik, P. G., Simpson, C., Tittensor, D. P., Byrnes, J. E., … Pandolfi, J. M.. (2015). Paleontological baselines for evaluating extinction risk in the modern oceans. *Science*, *348*, 567–570.

Gislason, H., Pope, J. G., Rice, J. C., & Daan, N. (2008). Coexistence in North Sea fish communities: Implications for growth and natural mortality. *ICES Journal of Marine Science*, *65*, 514–530.

Grueber, C. E., Nakagawa, S., Laws, R. J., & Jamieson, I. G. (2011). Multimodel inference in ecology and evolution: Challenges and solutions. *Journal of Evolutionary Biology*, *24*, 699–711.

Harsch, M. A., Zhou, Y., HilleRisLambers, J., & Kot, M. (2014). Keeping pace with climate change: Stage-structured moving-habitat models. *American Naturalist*, *184*, 25–37.

Hastie, T., Tibshirani, R., & Friedman, J. (2009). *The elements of statistical learning: Data mining, inference, and prediction*, 2nd edn. New York: Springer.

Hilborn, R. W., & Walters, C. (1992). *Quantitative fisheries stock assessment: Choice, dynamics, and uncertainty*. London, UK: Chapman and Hall.

Hutchings, J. A., Minto, C., Ricard, D., Baum, J. K., & Jensen, O. P. (2010). Trends in the abundance of marine fishes. *Canadian Journal of Fisheries and Aquatic Sciences*, *67*, 1205–1210.

IPCC (2013). *Climate Change 2013: The Physical Science Basis. Contribution of Working Group I to the Fifth Assessment Report of the Intergovernmental Panel on Climate Change*. Cambridge, UK and New York, NY: Cambridge University Press.

IUCN (2015). The IUCN Red List of Threatened Species. Version 2015.1. Technical report.

Johnson, K. F., Monnahan, C. C., McGilliard, C. R., Vert-pre, K. A., Anderson, S. C., Cunningham, C. J., … Punt, A. E. (2015). Time-varying natural mortality in fisheries stock assessment models: Identifying a default approach. *ICES Journal of Marine Science*, *72*, 137–150.

Johnson, J. B., & Omland, K. S. (2004). Model selection in ecology and evolution. *Trends in Ecology & Evolution*, *19*, 101–108.

Kell, L. T., Mosqueira, I., Grosjean, P., Fromentin, J.-M., Garcia, D., Hillary, R., … Scott, R. D. (2007). FLR: An open-source framework for the evaluation and development of management strategies. *ICES Journal of Marine Science*, *64*, 640–646.

Knutti, R., Furrer, R., Tebaldi, C., Cermak, J., & Meehl, G. A. (2009). Challenges in combining projections from multiple climate models. *Journal of Climate*, *23*, 2739–2758.

Krishnamurti, T. N., Kishtawal, C. M., LaRow, T. E., Bachiochi, D. R., Zhang, Z., Williford, C. E., … Surendran, S. (1999). Improved weather and seasonal climate forecasts from multimodel superensemble. *Science*, *285*, 1548–1550.

Liaw, A., & Wiener, M. (2002). Classification and regression by random forest. *R News*, *2*, 18–22.

Martell, S., & Froese, R. (2013). A simple method for estimating MSY from catch and resilience. *Fish and Fisheries*, *14*, 504–514.

Mote, P. W., Allen, M. R., Jones, R. G., Li, S., Mera, R., Rupp, D. E. … Vickers, D. (2015). Superensemble regional climate modeling for the western US. *Bulletin of the American Meteorological Society*, *97*, 203–215.

Murphy, J. M., Sexton, D. M. H., Barnett, D. N., Jones, G. S., Webb, M. J., Collins, M., & Stainforth, D. A. (2004). Quantification of modelling uncertainties in a large ensemble of climate change simulations. *Nature*, *430*, 768–772.

Pinsky, M. L., Jensen, O. P., Ricard, D., & Palumbi, S. R. (2011). Unexpected patterns of fisheries collapse in the world's oceans. *Proceedings of the National Academy of Sciences of the United States of America*, *108*, 8317–8322.

R Core Team (2015). *R: A language and environment for statistical computing*. Vienna, Austria: R Foundation for Statistical Computing.

Rassweiler, A., Costello, C., Hilborn, R., & Siegel, D. A. (2014). Integrating scientific guidance into marine spatial planning. *Proceedings of the Royal Society of London. Series B: Biological Sciences*, *281*, 20132252.

Ricard, D., Minto, C., Jensen, O. P., & Baum, J. K. (2012). Examining the knowledge base and status of commercially exploited marine species with the RAM Legacy Stock Assessment Database. *Fish and Fisheries*, *13*, 380–398.

Ridgeway, G. (2015) *gbm: generalized boosted regression models*. R package version 2.1.1.

Rosenberg, A. A., Fogarty, M. J., Cooper, A. B., Dickey-Collas, M., Fulton, E. A., Gutiérrez, N. L., … Ye, Y. (2014). *Developing new approaches to global stock status assessment and fishery production potential of the seas*. Rome, Italy: FAO Fisheries and Aquaculture Circular 1086.

Rykiel, J. Jr (1996). Testing ecological models: The meaning of validation. *Ecological Modelling*, *90*, 229–244.

Sampson, D. B., & Scott, R. D. (2012). An exploration of the shapes and stability of population–selection curves. *Fish and Fisheries*, *13*, 89–104.

Schaefer, M. B. (1954). Some aspects of the dynamics of populations important to the management of the commercial marine fisheries. *Inter-American Tropical Tuna Commission Bulletin*, *1*, 23–56.

Tebaldi, C., & Knutti, R. (2007). The use of the multi-model ensemble in probabilistic climate projections. *Philosophical Transactions of the Royal Society of London A: Mathematical, Physical and Engineering Sciences*, *365*, 2053–2075.

Theil, H. (1950). A rank-invariant method of linear and polynomial regression analysis, I, II, and III. *Nederlandsche Akad. van Wetenschappen Proc.*, *53*, 386–392, 521–525, and 1397–1412.

Thomson, M. C., Doblas-Reyes, F. J., Mason, S. J., Hagedorn, R., Connor, S. J., Phindela, T., … Palmer, T. N. (2006). Malaria early warnings based on seasonal climate forecasts from multi-model ensembles. *Nature*, *439*, 576–579.

Thorson, J. T., Minto, C., Minte-Vera, C. V., Kleisner, K. M., & Longo, C. (2013). A new role for effort dynamics in the theory of harvested populations and data-poor stock assessment. *Canadian Journal of Fisheries and Aquatic Sciences*, *70*, 1829–1844.

Thorson, J. T., Ono, K., & Munch, S. B. (2014). A Bayesian approach to identifying and compensating for model misspecification in population models. *Ecology*, *95*, 329–341.

Vasconcellos, M., & Cochrane, K. (2005). Overview of world status of data-limited fisheries: Inferences from landings statistics. In G. H. Kruse, V. F. Gallucci, D. E. Hay, R. I. Perry, R. M. Peterman, T. C. Shirley, … & D. Woodby (Eds.), *Fisheries assessment and management in data-limited situations* (pp. 1–20). Fairbanks: Alaska Sea Grant, University of Alaska.

Walther, B. A., & Moore, J. L. (2005). The concepts of bias, precision and accuracy, and their use in testing the performance of species richness estimators, with a literature review of estimator performance. *Ecography*, *28*, 815–829.

Weigel, A. P., Knutti, R., Liniger, M. A., & Appenzeller, C. (2010). Risks of model weighting in multimodel climate projections. *Journal of Climate*, *23*, 4175–4191.

Yeakel, J. D., Stiefs, D., Novak, M., & Gross, T. (2011). Generalized modeling of ecological population dynamics. *Journal of Theoretical Biology*, *4*, 179–194.

Yun, W. T., Stefanova, L., Mitra, A. K., Kumar, T. S. V. V., Dewar, W., & Krishnamurti, T. N. (2005). A multi-model superensemble algorithm for seasonal climate prediction using DEMETER forecasts. *Tellus A*, *57*, 280–289.

SUPPORTING INFORMATION

Additional Supporting Information may be found online in the supporting information tab for this article.

How to cite this article: Anderson SC, Cooper AB, Jensen OP, et al. Improving estimates of population status and trend with superensemble models. *Fish Fish*. 2017;18:732–741. https://doi.org/10.1111/faf.12200